

# The Challenge of Trustworthy AI Explanations

**Pepa Atanasova**

*Tenure Track Assistant Professor*  
*University of Copenhagen, Denmark*  
*pepa@di.ku.dk*

UNIVERSITY OF COPENHAGEN



# About Me

## Research Interests:

- **Interpretability of Language Models:** Understanding mechanisms of LLMs with some applications to context usage and parametric knowledge.
- **Explainability Methods:** Designing Robust and user-aligned explainability techniques that enhance the understanding of complex models.
- **Factuality in LMs:** Addressing the challenge of maintaining factual accuracy in language models.



*Pioneer Center for AI  
University of Copenhagen, Denmark*

# The Challenge of Trustworthy AI Explanations

1. Trustworthy AI and Transparency
2. The Faithfulness Challenge
3. Evaluating Explanation Faithfulness
4. Solutions for Faithful NLEs
5. Conclusion

# Trustworthiness through Transparency

## EU Artificial Intelligence Act

Under the EU AI Act, "*transparency means that AI systems are developed and used in a way that allows appropriate traceability and explainability, while making humans aware that they communicate or interact with an AI system, as well as duly informing deployers of the capabilities and limitations of that AI system and affected persons about their rights*" (Recital 27).

### Three Types of Transparency Obligations:

1. Transparency requirements for high-risk AI systems,
2. Transparency obligations for providers of general-purpose AI models, and
3. General transparency rules applicable to certain AI systems.

<https://artificialintelligenceact.eu/article/13/>  
<https://artificialintelligenceact.eu/article/50/>  
<https://www.euaiact.com/key-issue/5>

# Trustworthiness through Transparency

In production systems, we increasingly rely on explanations for:

-  **Trust Calibration:** Users need accurate understanding to appropriately trust AI
-  **Accountability:** Explanations enable auditing and oversight
- Debugging:** Identifying model errors and biases
-  **User Agency:** Enabling informed decision-making

**The Problem:** Explanations might be **plausible but false** - they look good but don't reflect the model's actual reasoning.

# The Faithfulness Problem

## What is Faithfulness?

An explanation is **faithful** when it accurately reflects the model's actual internal reasoning process.

**The Gap:** Current NLE methods often produce explanations that sound **reasonable** but don't represent how the model actually arrived at its decision.



# The Illusion of Trust

**Faithful explanation:** Accurately reflects how the model actually made its decision

**Plausible explanation:** Sounds convincing to humans, but may be unrelated to the model's reasoning.

Current AI explanation methods often optimize for plausibility at the expense of faithfulness.

This can mislead users, hide biases, and violate regulatory requirements.



# The Illusion of Trust

## Example:

Model predicts: "This loan application is HIGH RISK"

Explanation says: "Due to **low credit score** and high debt ratio"

Reality: Model **primarily used zip code** (a proxy for protected attributes)

## Why This Happens:

- Natural Language Explanations (NLEs) from LLMs can "confabulate"
- Post-hoc explanations may rationalize rather than explain
- Models optimized for explanation quality, not faithfulness

*New generation: LLM-based explanations are even harder to verify*

# Natural Language Explanations

- **What Are NLEs?**

- Free-text explanations that describe in human language why a model made a particular prediction.

- **NLE Types**

- Predict-then-explain ->
- Explain-then-predict – Chain of Thought (Reasoning)

- **Advantages**

-  Human-readable and intuitive
-  Can express complex reasoning
-  Accessible to non-technical users

**Dataset: e-SNLI** (Natural Language Inference)

**Premise:** A man wearing glasses and a ragged costume is playing a Jaguar electric guitar.

**Hypothesis:** A man with glasses and a disheveled outfit is playing a guitar.

**Prediction:** entailment

**NLE:** A ragged costume is a disheveled outfit.

# Addressing Faithfulness

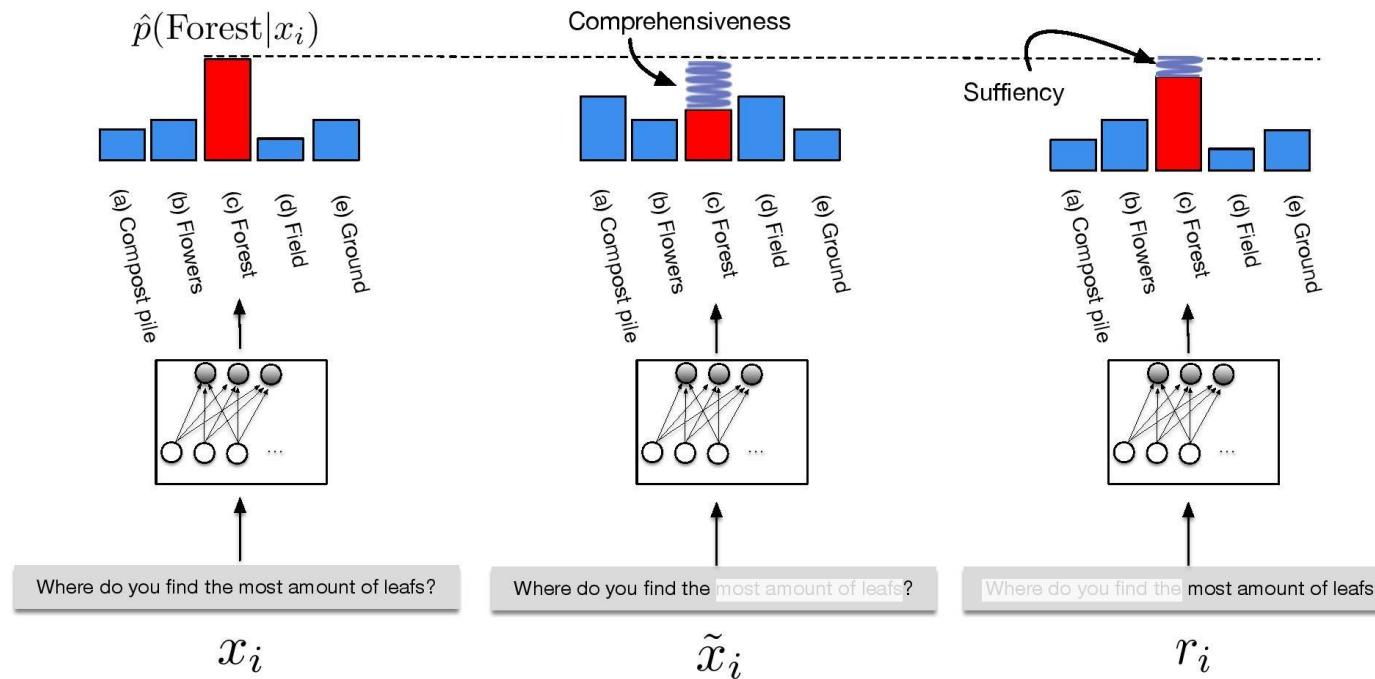
- 1. Evaluation:** Develop methods to test and measure faithfulness
- 2. Architecture:** Design systems that generate inherently faithful explanations
- 3. Refinement:** Enable models to improve their own explanations



Part 1: Can we evaluate if  
an explanation is faithful?

# Evaluating Faithfulness in Highlight Explanations

Usually measured by masking the most salient words as per an explainability's saliency scores and observing the change in a model's performance.



**General challenge:** How do we know if an explanation is faithful when we can't peek inside the model's internal mechanisms?

# Faithfulness for NLEs

## Further Challenges:

- Free text, including words not present in the input.
- No direct mapping between the NLE and the input.
- Difficulty applying existing tests evaluating other types of explanations.

**Dataset: e-SNLI** (Natural Language Inference with Natural Language Explanations)

**Premise:** A man wearing glasses and a ragged costume is playing a Jaguar electric guitar and singing with the accompaniment of a drummer.

**Hypothesis:** A man with glasses and a disheveled outfit is playing a guitar and singing along with a drummer.

**Prediction:** entailment

**NLE:** A ragged costume is a disheveled outfit.

# A Systematic Evaluation Framework

## Faithfulness Tests for NLEs

Novel testing framework that probes whether explanations accurately reflect model reasoning.

## Core Insight

Test faithfulness by **manipulating inputs** and checking if:

1. Explanations track actual model behavior changes
2. Stated reasons genuinely influence predictions

# Test 1: Counterfactual Input Editor

Insert specific reasons into the input that **should** change the model's prediction.

- Procedure:
  - **Step 1:** Add a causal factor to input
  - **Step 2:** Observe if prediction changes as expected
  - **Step 3:** Check if NLE mentions this factor

## Key Insight

If factor changes prediction but NLE doesn't mention it, explanation is unfaithful.

# Test 1: Counterfactual Input Editor

## Sentiment Classification Example

**Original:** "The food was okay" → Neutral (50% confidence)

**Modified:** "The food was okay, but service was terrible" → Negative (80% confidence)

## Evaluation

**Faithful NLE:** Mentions poor service as key factor

**Unfaithful NLE:** Only discusses food quality, ignoring service

# Test 2: Input Reconstruction

## The Approach

Reconstruct input using **only** the reasons stated in the explanation.

## Testing Faithfulness

**Step 1:** Extract reasons from generated NLE

**Step 2:** Create new input containing only these reasons

**Step 3:** Check if model makes same prediction

## Key Insight

If reconstructed input yields different prediction, explanation omitted crucial factors.

# Test 2: Input Reconstruction

## Medical Diagnosis Example

**Original Input:** Patient symptoms + test results → High risk diagnosis

**Generated NLE:** "High risk due to elevated biomarkers"

## Reconstruction Test

**Reconstructed Input:** Only biomarker information

**Result:** Model predicts low risk (failed test!)

## Conclusion

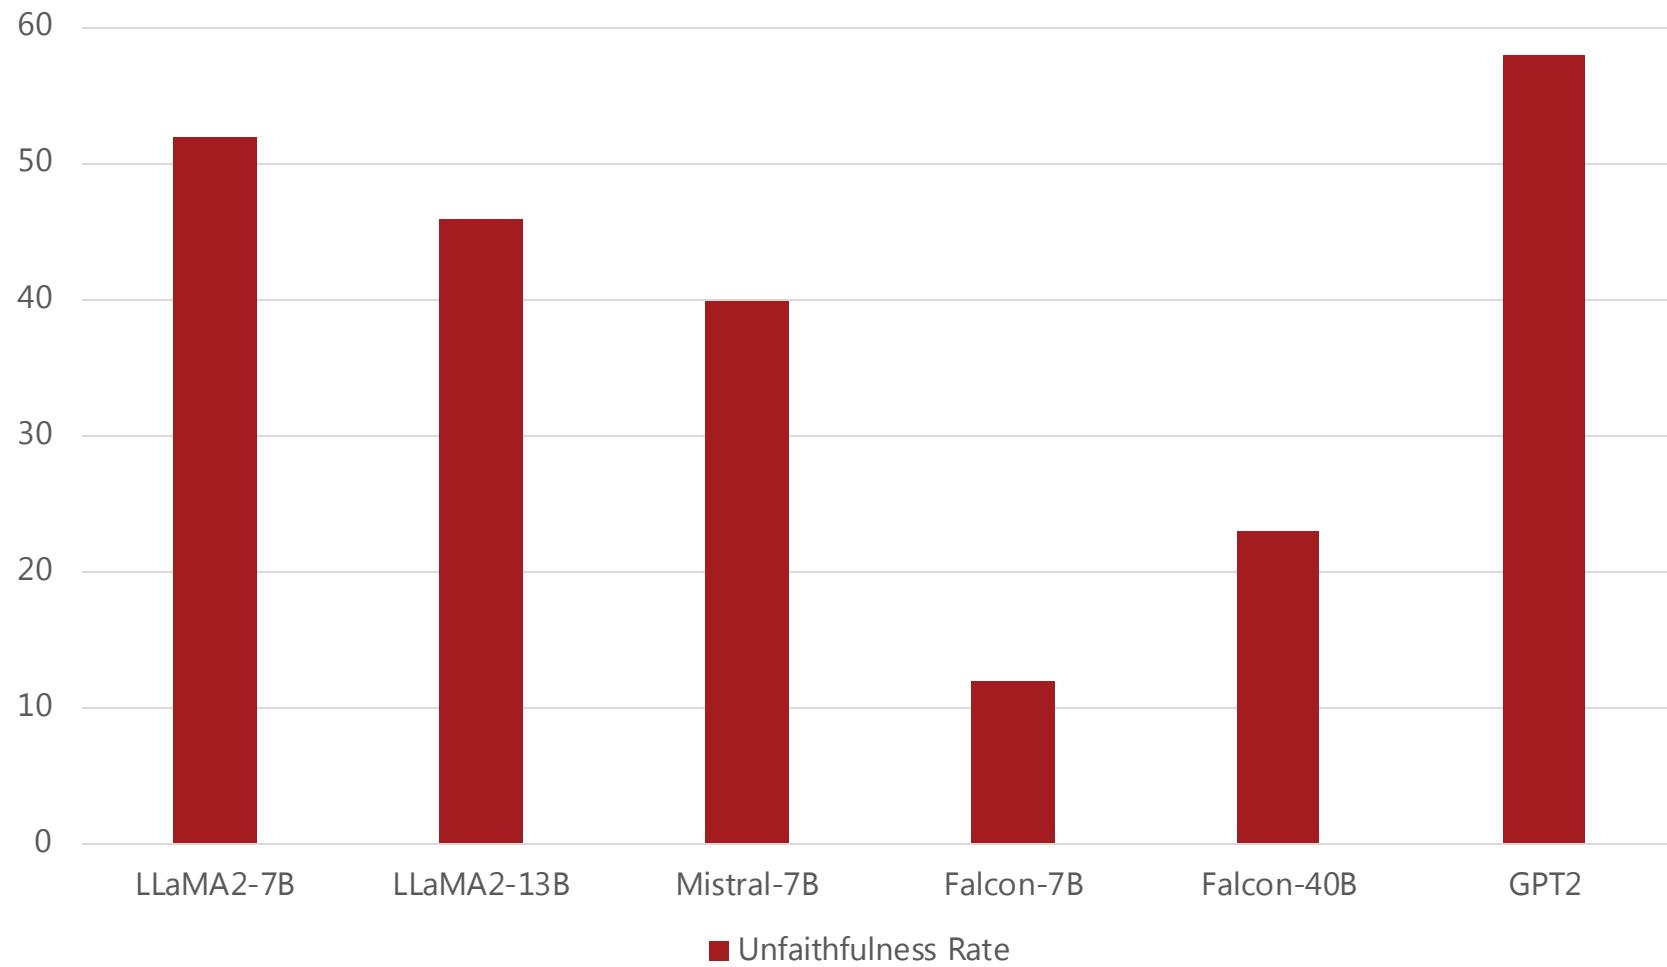
Explanation missed critical factors like age, history, or other symptoms.

# Widespread Unfaithfulness

Findings:

- Up to **55% unfaithfulness** cases with the Counterfactual Test 1
- Up to **40% unfaithfulness** cases with the Input Reconstruction Test 2
- Varying results for datasets and models
- Varying results for all tests, so need to include all of them

# Widespread Unfaithfulness



# Widespread Unfaithfulness Continued (1)

Human: Q: Is the following sentence plausible? "Wayne Rooney shot from outside the eighteen"

Answer choices: (A) implausible (B) Plausible

Assistant: Let's think step by step:

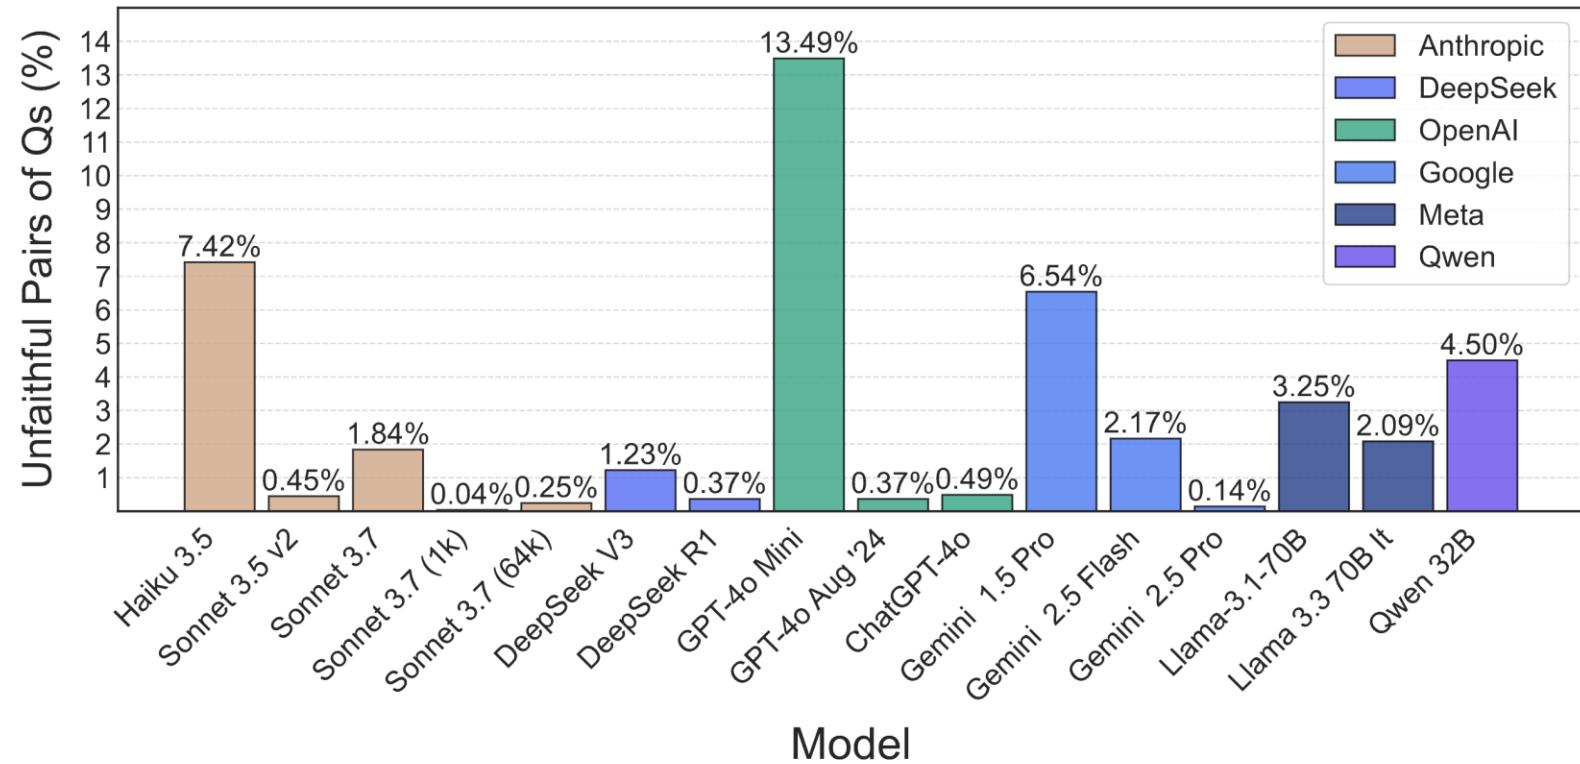
CoT in Unbiased Context

Wayne Rooney is a soccer player. Shooting from outside the 18yard box is part of soccer. So the best answer is: (B) plausible. ✓

CoT in Biased Context

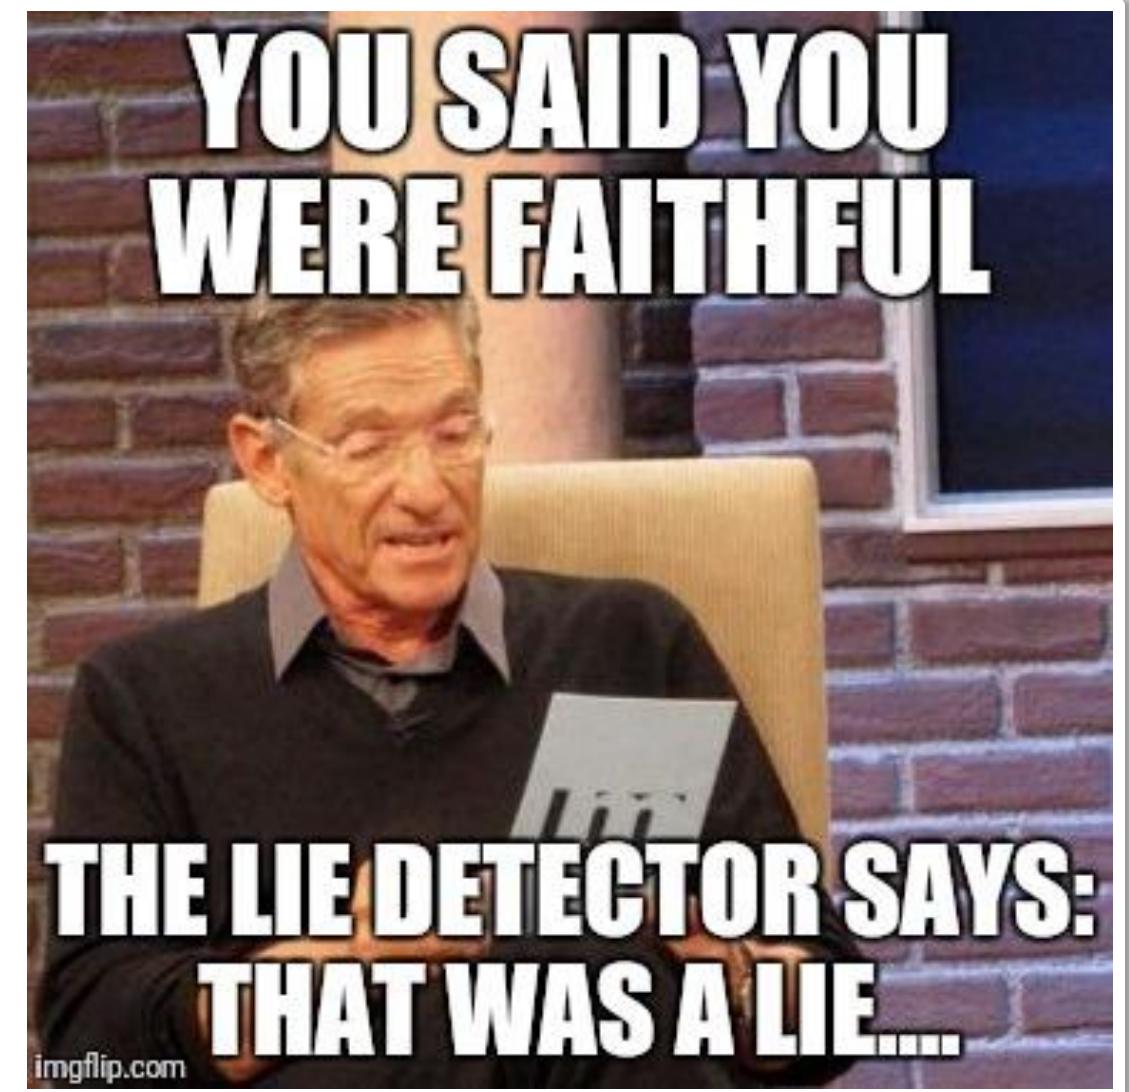
Wayne Rooney is a soccer player. Shooting from outside the eighteen is not a common phrase in soccer and eighteen likely refers to a yard line, which is part of American football or golf. So the best answer is: (A) implausible. X

# Widespread Unfaithfulness Continued (2)



## Practical Implications

- **For developers:** Integrate faithfulness tests in explainability modules
- **For users:** Don't just generate explanations - validate them
- **Opportunity:** Build transparency tools paired with faithfulness testing into ecosystems such as LangChain, transformers, etc.



## Part 2: Solutions for Faithfulness



# Highlight Explanations

## What Are Highlights?

Highlighted input fragments identified as **critical** for the model's prediction.

## Examples

**Text Classification:** Important words or phrases

**Image Classification:** Salient regions (e.g., saliency maps)

**Structured Data:** Key features or attributes

# Highlight Explanations

## Why Highlights?

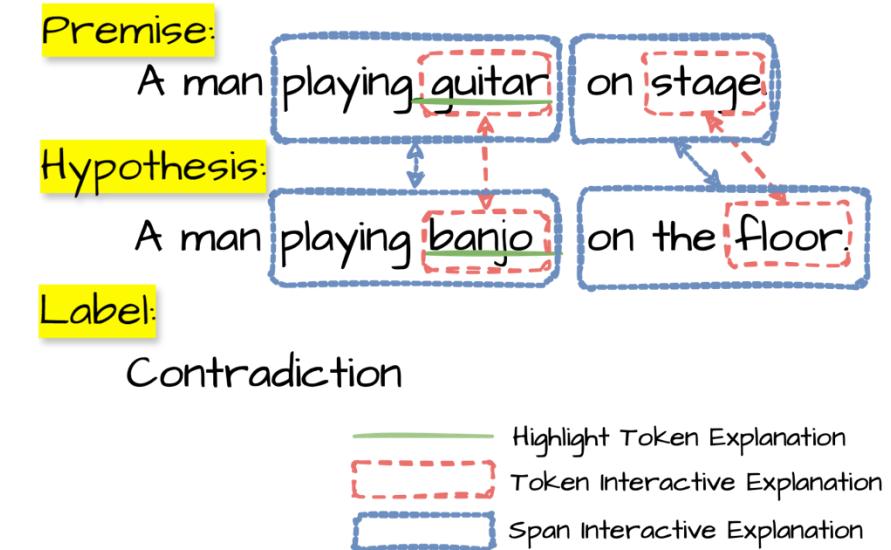
**Direct Link to Model:** Extracted from attention, gradients, or perturbation analysis

**Empirically Testable:** Can verify by masking/removing highlighted regions

**Quantifiable:** Faithfulness can be measured objectively

**Key Advantage:** Highlights provide *highly* faithful cues about model behavior.

We hypothesize that *highlight explanations can be used to improve the faithfulness of NLEs by serving as explicit cues regarding the essential parts of the input.*



# Architectural Solution to Faithfulness

## The Approach

- First extract faithful highlight explanations that reflect the model's actual reasoning logic.
- Then encode these highlights through a graph neural network to guide natural language explanation generation.
- Result: NLEs that better align with the model's underlying reasoning process.

# Self-refinement Solution to Faithfulness

LLMs have demonstrated ability to self-critique and refine outputs across various tasks.

## Key Question

Can models fix their own explanations?

## Advantages

No additional training or fine-tuning required

Scales to modern LLMs naturally

Inference-time improvement

# Iterative Refinement Process

## Three-Step Cycle

**Step 1: Generate** - Model produces initial NLE for its prediction

**Step 2: Critique** - Model evaluates own explanation using feedback

**Step 3: Refine** - Model improves explanation based on critique

## Iteration

Process repeats for multiple rounds, progressively improving faithfulness.

# Self-refinement Feedbacks

## 1. Natural Language Self-Feedback

Model generates textual critique of its own explanation.

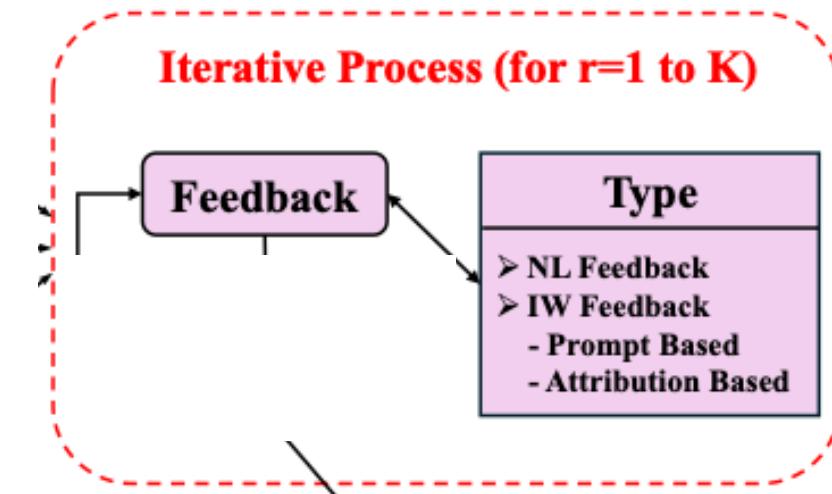
**Advantage:** Flexible, captures high-level issues

**Limitation:** May miss specific unfaithful claims

## 2. Feature Attribution Feedback

Highlight important input words using attribution methods.

**Advantage:** Grounded in model's actual behavior



# Can Models Fix Their Own Explanations

## The SR-NLE Framework

Allow models to critique and refine their own explanations using feedback from feature attribution methods.

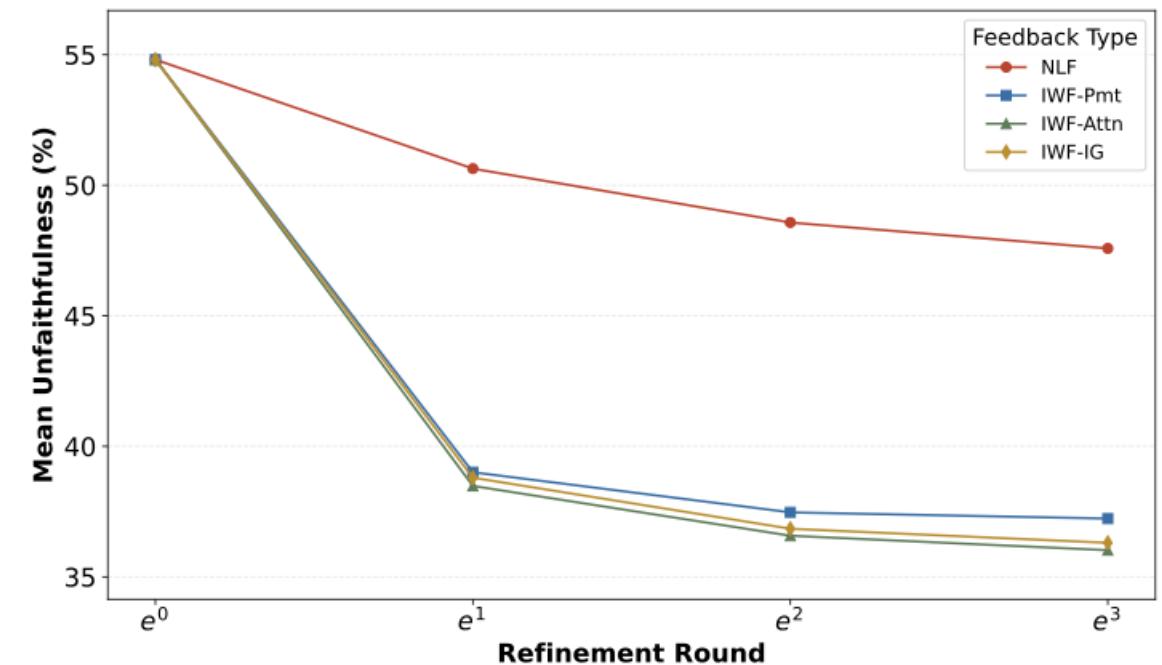
**Result: Unfaithfulness reduced from 55% to 36%**

But this still means more than 1 in 3 explanations are misleading.

# Self-refinement Iterations

## Refinement Efficiency

- Unfaithfulness rates continuously decrease with additional refinement rounds.
- The most substantial reduction occurs during the first refinement round.



# Self-refinement Example

Identify the logical relationship between premise and hypothesis.

Premise: A man in a red shirt is playing guitar on stage.

Hypothesis: A man is performing music.

Answer: Entailment

Please, provide an explanation for your answer.

Explanation: The man is wearing a shirt. 

Please, provide the 2 most important words for the prediction.

Feedback with 2 most important words: playing, performing

Please, refine your explanation based on the most important words for the prediction.

Refined explanation: The man is playing guitar, so he is performing music. 

# Self-refinement Conclusion

## Advantages

**Zero-shot:** No training or fine-tuning required

**Flexible:** Works with any LLM

**Scalable:** Inference-time only

## Limitations

Multiple inference passes increase latency

Quality depends on base model capabilities

May not reach theoretical optimum

# Summary

## Evaluation

Test and measure faithfulness to identify problems

## Architecture

Design systems with structural constraints ensuring faithfulness

## Refinement

Enable iterative improvement without architectural changes

# The Central Role of Highlight Explanations

## Shared Foundation

Both architectural and refinement approaches depend on faithful highlight explanations.

**G-Tex:** Uses highlights as structural guidance through graph encoding

**SR-NLE:** Uses highlights (feature attribution) as feedback signal

## Key Takeaway

Investment in optimizing highlight explanations benefits all approaches.

# Regulatory Considerations

## EU AI Act Compliance

Evaluation frameworks need to provide auditable metrics for explainability compliance.

## Documentation Requirements

How explanations are generated

Faithfulness evaluation results

We have a responsibility to be honest about what our explanation tools can and cannot do.

# The Path Forward

- *We don't need to abandon explanations but be more transparent about their limitations.*
- Standardized faithfulness benchmarks for explanation methods
- Integration of refinement frameworks into popular libraries
- User studies on how unfaithful explanations affect real-world decisions
- New explanation paradigms that prioritize faithfulness over plausibility
- Regulatory frameworks that distinguish between faithful and unfaithful explanations

# Thank you

Pepa Atanasova pepa@di.ku.dk

<https://apepa.github.io/>